

Latency and modality support of API-generated pseudo reports

Modality support rate

1.00

0.95

0.90

0.85

0.80

12

14

16

18

20

22

24

26

Mean latency per case (s)

GPT-5.5
(n=100)

Qwen-Plus
(n=10)

Gemini-3.1-Pro
(n=9)

GLM-4.7-Flash
(n=10)

Completion tokens
600
1,200
1,800

